

Mission1

October 31, 2024

```
[ ]: from google.colab import drive
import os, sys

#
drive.mount('/content/drive')

#
%cd /content/drive/MyDrive/

#
!mkdir -p Mission1
```

Mounted at /content/drive
/content/drive/MyDrive

1 Mission 1:

1.1 Mission 1-1

1. “{W/T}{ ID}{ }{ }{ }.jpg” . 2. & . 3. Train Validation

```
[ ]: import os
import pandas as pd
from IPython.display import display, display_html

# ( )
training_image_list = os.listdir('dataset/origin_dataset/training_image')
validation_image_list = os.listdir('dataset/origin_dataset/validation_image')

# train validation
print(f"train image: {len(training_image_list)}, validation image: {len(validation_image_list)}")

# -> Pandas
def path2data(image_list):
    data = []
    for i in range(len(image_list)):
```

```

        meta_data = image_list[i].split('_')

        row = {
            'file_name': image_list[i][: -4],      # ".jpg"
            'wt': meta_data[0],
            'image_id': meta_data[1],
            'time': meta_data[2],
            'style': meta_data[3],
            'gender': meta_data[4][0]
        }
        data.append(row)
    return pd.DataFrame(data)

#
training_df = path2data(training_image_list)
validation_df = path2data(validation_image_list)

#      HTML
def display_left(*args):
    html_str = ''
    for df in args:
        html_str += f'<div style="margin-right:30px;">{df.to_html()}</div>'
    display_html(f'<div style="display: flex;">{html_str}</div>', raw=True)

#
display_left(training_df.head(), validation_df.head())

#
training_df.to_csv('Mission1/training_df.csv', index=False)
validation_df.to_csv('Mission1/validation_df.csv', index=False)

```

train image: 4070, validation image: 951

```

[ ]: #
training_count_data = training_df[['gender', 'style', 'image_id']].
    ↳groupby(['gender', 'style']).count()
validation_count_data = validation_df[['gender', 'style', 'image_id']].
    ↳groupby(['gender', 'style']).count()

#
training_count_data.rename(columns={'image_id': 'training count'}, inplace=True)
validation_count_data.rename(columns={'image_id': 'validation count'},
    ↳inplace=True)

#
display_left(training_count_data, validation_count_data)

```

```
#
training_count_data.to_csv('Mission1/training_count_data.csv')
validation_count_data.to_csv('Mission1/validation_count_data.csv')
```

1.1.1

(1)

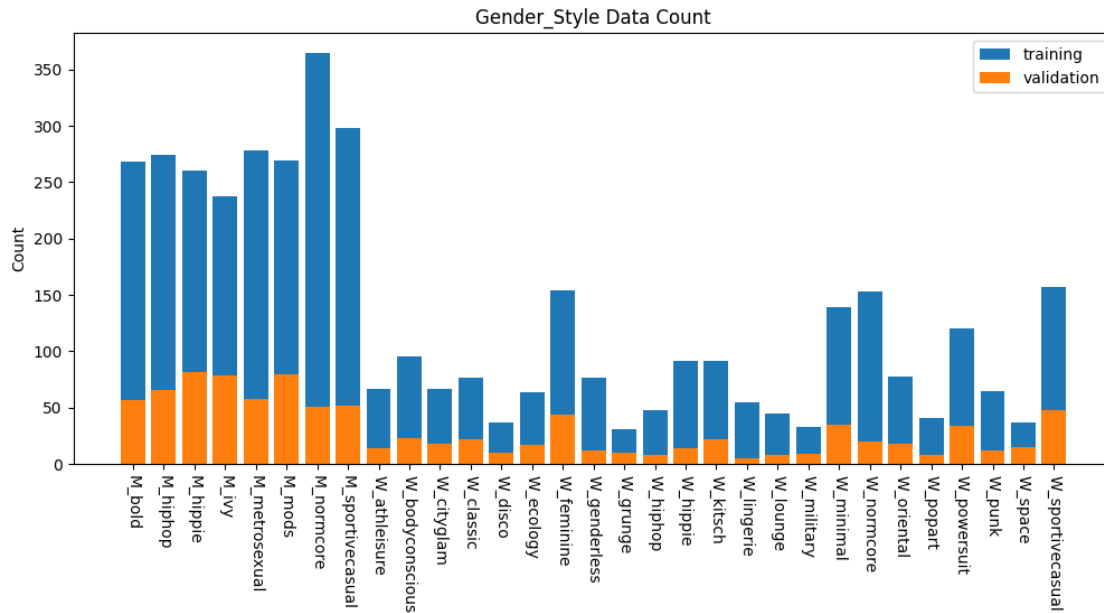
```
[ ]: import matplotlib.pyplot as plt

# Group
training_count_data_re = training_count_data.reset_index()
validation_count_data_re = validation_count_data.reset_index()

# Gender Style
training_count_data_re['name'] = training_count_data_re['gender'] + '_' +
    ↪training_count_data_re['style']
validation_count_data_re['name'] = validation_count_data_re['gender'] + '_' +
    ↪validation_count_data_re['style']

#
plt.figure(figsize=(12, 5))
plt.title('Gender_Style Data Count')
plt.ylabel('Count')
plt.xticks(rotation=270)
plt.bar(training_count_data_re['name'], training_count_data_re['training_
    ↪count'], label="training")
plt.bar(validation_count_data_re['name'], validation_count_data_re['validation_
    ↪count'], label="validation")

plt.legend()
plt.savefig('Mission1/gender_style_data_count.png', bbox_inches='tight') #
    ↪
plt.show()
```



(2)

```
[ ]: import matplotlib.pyplot as plt

#
training_df_tt = training_df.copy()
validation_df_tt = validation_df.copy()

#      00 -> 2000, 80 -> 1980
def time2year(time):
    g20 = ['00', '10', '19']
    if time in g20:
        return '20' + time
    else:
        return '19' + time


training_df_tt['time'] = training_df_tt['time'].apply(time2year)
validation_df_tt['time'] = validation_df_tt['time'].apply(time2year)

#
training_by_year = training_df_tt.groupby('time').count()
validation_by_year = validation_df_tt.groupby('time').count()

#
plt.figure(figsize=(12, 5))
plt.title('Year Data Count')
plt.ylabel('Count')
```

```
plt.bar(training_by_year.index, training_by_year['image_id'], label="training")
plt.bar(validation_by_year.index, validation_by_year['image_id'],
        label="validation")

plt.legend()
plt.savefig('Mission1/year_data_count.png', bbox_inches='tight') #
plt.show()
```



(3) ID

- ID ,

```
[ ]: import matplotlib.pyplot as plt

# image_id
training_duplicates = training_df[training_df.duplicated(['image_id'],
        keep=False)]
validation_duplicates = validation_df[validation_df.duplicated(['image_id'],
        keep=False)]

#
training_duplicates.sort_values(by=['image_id'])
validation_duplicates.sort_values(by=['image_id'])

#
display_left(training_duplicates, validation_duplicates)
```

1.2 Mission 1-2

1. “ResNet18 & . 2. parameters , pretrained weights . 3.

```
[ ]: #      import
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms, models
import random
import numpy as np

#
def set_random_seed(seed_value=42):
    # Python
    random.seed(seed_value)
    # NumPy
    np.random.seed(seed_value)
    # PyTorch (CPU)
    torch.manual_seed(seed_value)
    # PyTorch (GPU)
    torch.cuda.manual_seed(seed_value)
    torch.cuda.manual_seed_all(seed_value)

    # CuDNN
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

set_random_seed()
```

1.2.1 1.

1. pandas

2. JSON .

```
[ ]: #####
# Preprocessing Metadata
#####
import json
import pandas as pd
from IPython.display import display, display_html

#
training_image_list = os.listdir('dataset/origin_dataset/training_image')
validation_image_list = os.listdir('dataset/origin_dataset/validation_image')

#      -> Pandas
def path2data(image_list):
```

```

data = []
for i in range(len(image_list)):
    meta_data = image_list[i].split('_')

    wt = meta_data[0]
    image_id = meta_data[1]
    time = meta_data[2]
    style = meta_data[3]
    gender = meta_data[4][0]

    row = {
        'wt': wt,
        'image_id': image_id,
        'time': time,
        'style': style,
        'gender': gender,
        'image_name': image_list[i]
    }
    data.append(row)
return pd.DataFrame(data)

#
training_df = path2data(training_image_list)
validation_df = path2data(validation_image_list)

#
training_df['label_name'] = training_df['gender'] + '_' + training_df['style']
validation_df['label_name'] = validation_df['gender'] + '_' +
    ↪validation_df['style']

#      HTML
def display_left(*args):
    html_str = ''
    for df in args:
        html_str += f'<div style="margin-right:30px;">{df.to_html()}</div>'
    display_html(f'<div style="display: flex;">{html_str}</div>', raw=True)

#
label_map = {}
for idx, label in enumerate(training_df['label_name'].unique()):
    label_map[label] = idx
    label_map[idx] = label

#
training_df['label'] = training_df['label_name'].map(label_map)

```

```

validation_df['label'] = validation_df['label_name'].map(label_map)

#
display_left(training_df.head(), validation_df.head())
print(len(training_df), len(validation_df))

#
training_df.to_csv('dataset/training_image_labels.csv', index=False)
validation_df.to_csv('dataset/validation_image_labels.csv', index=False)

#
with open('dataset/label_map.json', 'w') as f:
    json.dump(label_map, f)

#
print(f"{label_map[0]}, {label_map['W_minimal']}")

```

4070 951
W_minimal, 0

1.2.2 2.

1. ReSize .
2. Colab Google Drive I/O , Multi Thread .
3. I/O Tensor .

```

[ ]: #####
# Preprocessing File
#####

import os
import torch
from torchvision import transforms
from tqdm.notebook import tqdm
from PIL import Image
import pandas as pd
from concurrent.futures import ThreadPoolExecutor

#
def process_single_image(args):
    image_name, label, image_dir, transform = args
    image_path = os.path.join(image_dir, image_name)
    try:
        #
        image = Image.open(image_path).convert('RGB')
        #
        image = transform(image)
        return image, label

```



```

except Exception as e:
    print(f"Error processing {image_name}: {e}")
    return None

#
def process_images_parallel(metadata_df, image_dir, transform, cpu):
    inputs = [(row['image_name'], row['label'], image_dir, transform) for idx,
    ↪row in metadata_df.iterrows()]
    images = []
    labels = []

    # Thread Pool
    with ThreadPoolExecutor(max_workers=min(cpu, os.cpu_count())) as executor:
        # executor
        results = list(tqdm(executor.map(process_single_image, inputs),
    ↪total=len(inputs)))

    for result in results:
        if result is not None:
            image, label = result
            images.append(image)
            labels.append(label)

    #
    if not images:
        raise ValueError("No images were processed successfully.")

    images = torch.stack(images)
    labels = torch.tensor(labels, dtype=torch.long)
    return images, labels

#
def preprocess_image_to_tensor_pipeline(train_image_dir, val_image_dir,
    ↪tensor_dir, train_transform, test_transform, cpu=8):
    #
    os.makedirs(tensor_dir, exist_ok=True)

    #
    training_metadata_df = pd.read_csv('dataset/training_image_labels.csv')
    validation_metadata_df = pd.read_csv('dataset/validation_image_labels.csv')

    #
    print("Train Data Preprocessing")
    train_images, train_labels = process_images_parallel(training_metadata_df,
    ↪train_image_dir, train_transform, cpu=cpu)

```

```

    torch.save((train_images, train_labels), os.path.join(tensor_dir, 'train.
    pt'))

    #
    print("Validation Data Preprocessing")
    val_images, val_labels = process_images_parallel(validation_metadata_df,
    val_image_dir, test_transform, cpu=cpu)
    torch.save((val_images, val_labels), os.path.join(tensor_dir, 'val.pt'))

    print("Preprocessing Complete")

```

```

[ ]: # ( L4 GPU . 12 )
os.cpu_count()

```

[]: 12

```

[ ]: transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor()
])

!mkdir -p dataset/tensor

# (L4 8)
preprocess_image_to_tensor_pipeline(
    #
    train_image_dir = 'dataset/origin_dataset/training_image',
    val_image_dir = 'dataset/origin_dataset/validation_image',
    #
    tensor_dir = 'dataset/tensor/resize_tensor',
    #
    train_transform = transform,
    test_transform = transform,
    # ( .)
    cpu=8
)

```

Train Data Preprocessing

0%| | 0/4070 [00:00<?, ?it/s]

Validation Data Preprocessing

0%| | 0/951 [00:00<?, ?it/s]

Preprocessing Complete

1.2.3 3.

1. epoch .

```
[ ]: # GPU
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda:0

```
[ ]: #####
# Data Load
#####

from torch.utils.data import DataLoader
from torchvision import datasets, transforms, models

#
class CustomDataset(torch.utils.data.Dataset):
    def __init__(self, images, labels, transform=None):
        self.images = images
        self.labels = labels
        self.transform = transform

    def __getitem__(self, index):
        image = self.images[index]
        label = self.labels[index]

        if self.transform:
            image = self.transform(image)

        return image, label

    def __len__(self):
        return len(self.images)

def data_loading(tensor_dir, batch_size, train_transform, val_transform):
    """
    tensor_dir:
    batch_size:
    train_transform:
    val_transform:
    """

    #
    train_images, train_labels = torch.load(os.path.join(tensor_dir, 'train.
    ↪pt'), weights_only=False)
    val_images, val_labels = torch.load(os.path.join(tensor_dir, 'val.pt'),
    ↪weights_only=False)

    #
```

```

train_dataset = CustomDataset(train_images, train_labels,
↪transform=train_transform)
val_dataset = CustomDataset(val_images, val_labels, transform=val_transform)

#
train_loader = DataLoader(train_dataset, batch_size=batch_size,
↪shuffle=True, pin_memory=True) # [3] Pin-Memory
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False,
↪pin_memory=True)

#
print(f"Train Dataset: {len(train_dataset)}, Val Dataset:
↪{len(val_dataset)}")

return train_loader, val_loader

```

1.2.4 4.

1. .
2. Overfitting, Local Minimum .
3. .

```

[ ]: #####
# Training
#####

import matplotlib.pyplot as plt
from tqdm.notebook import tqdm, trange

def training(experiment_name, save_dir, data_loader, optimizer, criterion,
↪scheduler=None):
    """
    experiment_name:
    save_dir:
    data_loader:
    optimizer:
    criterion:
    batch_size:
    scheduler:
    """

    train_loader, val_loader = data_loader

    os.makedirs(save_dir, exist_ok=True)

```

```

os.makedirs(os.path.join(save_dir, experiment_name), exist_ok=True)

train_loss_chart = []
val_loss_chart = []

train_acc_chart = []
val_acc_chart = []

for epoch in trange(NUM_EPOCHS):
    model.train()

    train_loss = 0
    train_correct = 0
    total = 0
    for images, labels in tqdm(train_loader, leave=False):
        images, labels = images.to(device), labels.to(device)
        images, labels = images.to(device, non_blocking=True), labels.
↪to(device, non_blocking=True) # [4]

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        train_correct += (predicted == labels).sum().item()

    train_loss /= len(train_loader)
    train_acc = 100 * train_correct / total

    train_loss_chart.append(train_loss)
    train_acc_chart.append(train_acc)

    print(f"Epoch [{epoch+1}/{NUM_EPOCHS}], Train Loss: {train_loss:.4f}, ↪
↪Train Acc: {train_acc:.2f}%")

    model.eval()
    val_loss = 0
    val_correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in tqdm(val_loader, leave=False):

```

```

        images, labels = images.to(device), labels.to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)
        val_loss += loss.item()

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        val_correct += (predicted == labels).sum().item()

    val_loss /= len(val_loader)
    val_acc = 100 * val_correct / total

    # BestModel
    if epoch == 0 or val_acc > max(val_acc_chart):
        if experiment_name == 'best_acc_exp':
            torch.save(model.state_dict(), f'{save_dir}/{experiment_name}/
↪best_acc_model.pth')
        elif experiment_name == 'best_loss_exp':
            torch.save(model.state_dict(), f'{save_dir}/{experiment_name}/
↪best_loss_model.pth')
        else:
            torch.save(model.state_dict(), f'{save_dir}/{experiment_name}/
↪best_model.pth')

    val_loss_chart.append(val_loss)
    val_acc_chart.append(val_acc)

    print(f"Epoch [{epoch+1}/{NUM_EPOCHS}], Val Loss: {val_loss:.4f}, Val_
↪Acc: {val_acc:.2f}%")

    #
    if scheduler is not None:
        scheduler.step()

    #
    if epoch+1 >= 20 and (epoch+1) % 5 == 0:
        torch.save(model.state_dict(), f'{save_dir}/{experiment_name}/
↪model_{epoch+1}.pth')
        print(f"Model saved at epoch -> {epoch+1}")

    # loss, acc
    plt.figure(figsize=(10, 5))
    plt.subplot(1, 2, 1)
    plt.plot(train_loss_chart, label='Train Loss')
    plt.plot(val_loss_chart, label='Val Loss')
    plt.xlabel('Epoch')

```

```

plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss')
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(train_acc_chart, label='Train Accuracy')
plt.plot(val_acc_chart, label='Val Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')
plt.grid(True)

plt.tight_layout()

#
plt.savefig(f'{save_dir}/{experiment_name}/train_visualization.png')
plt.show()

#
display(pd.DataFrame({
    'epoch': range(1, NUM_EPOCHS+1),
    'Train Acc': train_acc_chart, 'Train Loss': train_loss_chart,
    'Val Acc': val_acc_chart, 'Val Loss': val_loss_chart
}))

#
torch.save(model.state_dict(), f'{save_dir}/{experiment_name}/
model_final({epoch+1}).pth')
print(f"Model saved at epoch {epoch+1}")

```

1.2.5 4.

1. , .
2. ACC .
3. 3-2 Loss .

1. Best Acc (Mission1-2) Best ACC

```

[ ]: #
LR = 0.001
BATCH_SIZE = 64
NUM_EPOCHS = 100
WEIGHT_DECAY = 5e-4

# ,
EXP_NAME = 'best_acc_exp'

```

```

SAVE_DIR = "Mission1"

#
train_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
val_transform = transforms.Compose([
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
data_loader = data_loading(
    tensor_dir='dataset/tensor/resize_tensor',
    batch_size=BATCH_SIZE,
    train_transform=train_transform,
    val_transform=val_transform
)

# ResNet18 ( X)
model = models.resnet18(weights=None, num_classes=31)
model = model.to(device)

#
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=LR, weight_decay=WEIGHT_DECAY)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.8)

#
training(
    experiment_name=EXP_NAME,
    save_dir=SAVE_DIR,
    data_loader=data_loader,
    optimizer=optimizer,
    criterion=criterion,
    scheduler=scheduler
)

```

Train Dataset: 4070, Val Dataset: 951

0%| | 0/100 [00:00<?, ?it/s]

0%| | 0/64 [00:00<?, ?it/s]

Epoch [1/100], Train Loss: 3.2918, Train Acc: 8.08%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [1/100], Val Loss: 3.2451, Val Acc: 8.62%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [2/100], Train Loss: 3.1859, Train Acc: 9.34%


```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [2/100], Val Loss: 3.1987, Val Acc: 8.62%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [3/100], Train Loss: 3.1537, Train Acc: 9.31%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [3/100], Val Loss: 3.2258, Val Acc: 7.57%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [4/100], Train Loss: 3.1275, Train Acc: 9.07%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [4/100], Val Loss: 3.1117, Val Acc: 8.94%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [5/100], Train Loss: 3.1097, Train Acc: 10.44%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [5/100], Val Loss: 3.1075, Val Acc: 10.30%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [6/100], Train Loss: 3.0818, Train Acc: 10.91%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [6/100], Val Loss: 3.1590, Val Acc: 8.62%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [7/100], Train Loss: 3.0633, Train Acc: 11.13%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [7/100], Val Loss: 3.2435, Val Acc: 8.73%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [8/100], Train Loss: 3.0566, Train Acc: 11.15%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [8/100], Val Loss: 3.0601, Val Acc: 10.83%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [9/100], Train Loss: 3.0382, Train Acc: 11.87%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [9/100], Val Loss: 3.0709, Val Acc: 9.04%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [10/100], Train Loss: 3.0268, Train Acc: 11.47%

```

```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [10/100], Val Loss: 3.1901, Val Acc: 7.57%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [11/100], Train Loss: 2.9921, Train Acc: 12.83%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [11/100], Val Loss: 3.0055, Val Acc: 11.99%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [12/100], Train Loss: 2.9513, Train Acc: 13.69%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [12/100], Val Loss: 3.0223, Val Acc: 12.93%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [13/100], Train Loss: 2.9335, Train Acc: 14.08%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [13/100], Val Loss: 2.9870, Val Acc: 14.83%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [14/100], Train Loss: 2.8982, Train Acc: 14.99%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [14/100], Val Loss: 3.0593, Val Acc: 11.99%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [15/100], Train Loss: 2.8877, Train Acc: 15.33%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [15/100], Val Loss: 3.0125, Val Acc: 11.88%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [16/100], Train Loss: 2.8557, Train Acc: 14.79%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [16/100], Val Loss: 3.5437, Val Acc: 9.99%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [17/100], Train Loss: 2.8294, Train Acc: 15.95%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [17/100], Val Loss: 3.0565, Val Acc: 11.46%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [18/100], Train Loss: 2.7857, Train Acc: 16.93%

```

```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [18/100], Val Loss: 2.8335, Val Acc: 15.67%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [19/100], Train Loss: 2.7570, Train Acc: 17.71%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [19/100], Val Loss: 2.8939, Val Acc: 14.30%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [20/100], Train Loss: 2.7028, Train Acc: 18.33%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [20/100], Val Loss: 3.0122, Val Acc: 13.77%
Model saved at epoch -> 20
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [21/100], Train Loss: 2.6222, Train Acc: 20.54%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [21/100], Val Loss: 2.7618, Val Acc: 17.46%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [22/100], Train Loss: 2.5600, Train Acc: 21.72%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [22/100], Val Loss: 2.7890, Val Acc: 16.61%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [23/100], Train Loss: 2.5162, Train Acc: 23.24%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [23/100], Val Loss: 2.7798, Val Acc: 18.82%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [24/100], Train Loss: 2.4413, Train Acc: 24.69%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [24/100], Val Loss: 2.6720, Val Acc: 21.77%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [25/100], Train Loss: 2.3924, Train Acc: 26.58%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [25/100], Val Loss: 3.0593, Val Acc: 15.46%
Model saved at epoch -> 25
0%|          | 0/64 [00:00<?, ?it/s]

```

Epoch [26/100], Train Loss: 2.2984, Train Acc: 29.12%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [26/100], Val Loss: 2.7446, Val Acc: 21.14%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [27/100], Train Loss: 2.1848, Train Acc: 32.46%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [27/100], Val Loss: 2.8954, Val Acc: 18.61%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [28/100], Train Loss: 2.0756, Train Acc: 36.63%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [28/100], Val Loss: 2.5430, Val Acc: 24.40%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [29/100], Train Loss: 1.9443, Train Acc: 38.62%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [29/100], Val Loss: 2.4288, Val Acc: 29.76%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [30/100], Train Loss: 1.7524, Train Acc: 45.87%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [30/100], Val Loss: 2.6128, Val Acc: 33.54%
Model saved at epoch -> 30
0%| | 0/64 [00:00<?, ?it/s]

Epoch [31/100], Train Loss: 1.4895, Train Acc: 56.09%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [31/100], Val Loss: 2.1838, Val Acc: 41.96%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [32/100], Train Loss: 1.1900, Train Acc: 65.65%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [32/100], Val Loss: 2.5570, Val Acc: 38.59%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [33/100], Train Loss: 1.0152, Train Acc: 70.34%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [33/100], Val Loss: 2.4625, Val Acc: 40.90%

```

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [34/100], Train Loss: 0.8674, Train Acc: 75.36%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [34/100], Val Loss: 2.0145, Val Acc: 50.68%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [35/100], Train Loss: 0.6065, Train Acc: 84.28%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [35/100], Val Loss: 2.2470, Val Acc: 51.10%
Model saved at epoch -> 35
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [36/100], Train Loss: 0.4830, Train Acc: 87.49%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [36/100], Val Loss: 2.0714, Val Acc: 54.89%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [37/100], Train Loss: 0.3699, Train Acc: 91.99%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [37/100], Val Loss: 2.2142, Val Acc: 55.31%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [38/100], Train Loss: 0.3006, Train Acc: 93.22%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [38/100], Val Loss: 1.9557, Val Acc: 59.52%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [39/100], Train Loss: 0.2377, Train Acc: 95.33%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [39/100], Val Loss: 1.9628, Val Acc: 60.04%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [40/100], Train Loss: 0.1728, Train Acc: 96.90%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [40/100], Val Loss: 2.0490, Val Acc: 60.04%
Model saved at epoch -> 40
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [41/100], Train Loss: 0.1331, Train Acc: 98.11%
0%|          | 0/15 [00:00<?, ?it/s]

```

Epoch [41/100], Val Loss: 1.8955, Val Acc: 62.88%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [42/100], Train Loss: 0.0857, Train Acc: 98.94%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [42/100], Val Loss: 1.9432, Val Acc: 62.36%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [43/100], Train Loss: 0.0626, Train Acc: 99.29%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [43/100], Val Loss: 1.8554, Val Acc: 62.99%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [44/100], Train Loss: 0.0509, Train Acc: 99.29%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [44/100], Val Loss: 1.8924, Val Acc: 62.04%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [45/100], Train Loss: 0.0362, Train Acc: 99.63%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [45/100], Val Loss: 1.8690, Val Acc: 62.78%

Model saved at epoch -> 45

0%| | 0/64 [00:00<?, ?it/s]

Epoch [46/100], Train Loss: 0.0412, Train Acc: 99.43%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [46/100], Val Loss: 1.9843, Val Acc: 63.09%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [47/100], Train Loss: 0.0379, Train Acc: 99.46%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [47/100], Val Loss: 1.8964, Val Acc: 62.88%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [48/100], Train Loss: 0.0295, Train Acc: 99.66%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [48/100], Val Loss: 1.8491, Val Acc: 63.09%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [49/100], Train Loss: 0.0352, Train Acc: 99.43%

```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [49/100], Val Loss: 1.9199, Val Acc: 62.36%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [50/100], Train Loss: 0.0230, Train Acc: 99.73%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [50/100], Val Loss: 1.8550, Val Acc: 63.51%
Model saved at epoch -> 50
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [51/100], Train Loss: 0.0208, Train Acc: 99.73%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [51/100], Val Loss: 1.8992, Val Acc: 62.99%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [52/100], Train Loss: 0.0225, Train Acc: 99.68%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [52/100], Val Loss: 1.8359, Val Acc: 62.78%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [53/100], Train Loss: 0.0141, Train Acc: 99.85%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [53/100], Val Loss: 1.8487, Val Acc: 63.83%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [54/100], Train Loss: 0.0181, Train Acc: 99.73%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [54/100], Val Loss: 1.9157, Val Acc: 63.51%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [55/100], Train Loss: 0.0162, Train Acc: 99.75%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [55/100], Val Loss: 1.8973, Val Acc: 62.78%
Model saved at epoch -> 55
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [56/100], Train Loss: 0.0105, Train Acc: 99.88%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [56/100], Val Loss: 1.8523, Val Acc: 63.72%
0%|          | 0/64 [00:00<?, ?it/s]

```

Epoch [57/100], Train Loss: 0.0166, Train Acc: 99.83%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [57/100], Val Loss: 2.0611, Val Acc: 63.62%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [58/100], Train Loss: 0.0316, Train Acc: 99.58%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [58/100], Val Loss: 1.9757, Val Acc: 63.30%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [59/100], Train Loss: 0.0254, Train Acc: 99.73%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [59/100], Val Loss: 1.8608, Val Acc: 62.99%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [60/100], Train Loss: 0.0571, Train Acc: 99.12%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [60/100], Val Loss: 2.2762, Val Acc: 57.62%
 Model saved at epoch -> 60

0%| | 0/64 [00:00<?, ?it/s]

Epoch [61/100], Train Loss: 0.2186, Train Acc: 95.18%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [61/100], Val Loss: 2.4833, Val Acc: 55.31%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [62/100], Train Loss: 0.2722, Train Acc: 93.00%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [62/100], Val Loss: 2.1811, Val Acc: 58.25%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [63/100], Train Loss: 0.2053, Train Acc: 95.14%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [63/100], Val Loss: 2.0399, Val Acc: 61.62%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [64/100], Train Loss: 0.1115, Train Acc: 97.84%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [64/100], Val Loss: 2.0098, Val Acc: 62.25%


```

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [65/100], Train Loss: 0.0646, Train Acc: 98.97%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [65/100], Val Loss: 1.9872, Val Acc: 62.04%
Model saved at epoch -> 65
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [66/100], Train Loss: 0.0411, Train Acc: 99.41%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [66/100], Val Loss: 1.9331, Val Acc: 63.72%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [67/100], Train Loss: 0.0244, Train Acc: 99.83%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [67/100], Val Loss: 1.8827, Val Acc: 63.20%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [68/100], Train Loss: 0.0257, Train Acc: 99.61%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [68/100], Val Loss: 1.9500, Val Acc: 62.36%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [69/100], Train Loss: 0.0157, Train Acc: 99.80%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [69/100], Val Loss: 1.8764, Val Acc: 63.62%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [70/100], Train Loss: 0.0109, Train Acc: 99.83%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [70/100], Val Loss: 1.9357, Val Acc: 63.20%
Model saved at epoch -> 70
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [71/100], Train Loss: 0.0137, Train Acc: 99.85%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [71/100], Val Loss: 1.8667, Val Acc: 63.41%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [72/100], Train Loss: 0.0092, Train Acc: 99.88%
0%|          | 0/15 [00:00<?, ?it/s]

```

Epoch [72/100], Val Loss: 1.8557, Val Acc: 63.83%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [73/100], Train Loss: 0.0103, Train Acc: 99.83%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [73/100], Val Loss: 1.8857, Val Acc: 62.99%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [74/100], Train Loss: 0.0090, Train Acc: 99.78%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [74/100], Val Loss: 1.8600, Val Acc: 63.20%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [75/100], Train Loss: 0.0090, Train Acc: 99.83%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [75/100], Val Loss: 1.8760, Val Acc: 63.41%

Model saved at epoch -> 75

0%| | 0/64 [00:00<?, ?it/s]

Epoch [76/100], Train Loss: 0.0103, Train Acc: 99.80%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [76/100], Val Loss: 1.8436, Val Acc: 63.20%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [77/100], Train Loss: 0.0065, Train Acc: 99.90%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [77/100], Val Loss: 1.8139, Val Acc: 63.51%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [78/100], Train Loss: 0.0049, Train Acc: 99.95%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [78/100], Val Loss: 1.8206, Val Acc: 63.51%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [79/100], Train Loss: 0.0048, Train Acc: 99.95%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [79/100], Val Loss: 1.8322, Val Acc: 63.20%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [80/100], Train Loss: 0.0055, Train Acc: 99.93%

```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [80/100], Val Loss: 1.7725, Val Acc: 63.41%
Model saved at epoch -> 80

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [81/100], Train Loss: 0.0067, Train Acc: 99.90%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [81/100], Val Loss: 1.8242, Val Acc: 63.41%

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [82/100], Train Loss: 0.0072, Train Acc: 99.88%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [82/100], Val Loss: 1.7816, Val Acc: 63.41%

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [83/100], Train Loss: 0.0047, Train Acc: 99.93%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [83/100], Val Loss: 1.7683, Val Acc: 63.41%

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [84/100], Train Loss: 0.0048, Train Acc: 99.93%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [84/100], Val Loss: 1.7553, Val Acc: 63.72%

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [85/100], Train Loss: 0.0052, Train Acc: 99.90%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [85/100], Val Loss: 1.7662, Val Acc: 63.41%
Model saved at epoch -> 85

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [86/100], Train Loss: 0.0077, Train Acc: 99.85%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [86/100], Val Loss: 1.7895, Val Acc: 63.72%

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [87/100], Train Loss: 0.0064, Train Acc: 99.90%

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [87/100], Val Loss: 1.7592, Val Acc: 63.20%

0%|          | 0/64 [00:00<?, ?it/s]

```

Epoch [88/100], Train Loss: 0.0074, Train Acc: 99.90%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [88/100], Val Loss: 1.9057, Val Acc: 63.41%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [89/100], Train Loss: 0.0057, Train Acc: 99.98%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [89/100], Val Loss: 1.7421, Val Acc: 63.41%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [90/100], Train Loss: 0.0045, Train Acc: 99.95%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [90/100], Val Loss: 1.7699, Val Acc: 63.72%
 Model saved at epoch -> 90

0%| | 0/64 [00:00<?, ?it/s]

Epoch [91/100], Train Loss: 0.0044, Train Acc: 99.98%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [91/100], Val Loss: 1.7647, Val Acc: 63.09%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [92/100], Train Loss: 0.0051, Train Acc: 99.98%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [92/100], Val Loss: 1.7200, Val Acc: 63.62%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [93/100], Train Loss: 0.0071, Train Acc: 99.93%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [93/100], Val Loss: 1.7834, Val Acc: 63.41%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [94/100], Train Loss: 0.0131, Train Acc: 99.85%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [94/100], Val Loss: 1.7928, Val Acc: 62.78%
 0%| | 0/64 [00:00<?, ?it/s]


Epoch [95/100], Train Loss: 0.0119, Train Acc: 99.93%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [95/100], Val Loss: 1.7267, Val Acc: 63.51%
 Model saved at epoch -> 95

```

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [96/100], Train Loss: 0.0058, Train Acc: 99.93%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [96/100], Val Loss: 1.7158, Val Acc: 63.20%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [97/100], Train Loss: 0.0096, Train Acc: 99.88%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [97/100], Val Loss: 1.8301, Val Acc: 63.09%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [98/100], Train Loss: 0.0117, Train Acc: 99.90%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [98/100], Val Loss: 1.7184, Val Acc: 62.99%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [99/100], Train Loss: 0.0078, Train Acc: 99.90%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [99/100], Val Loss: 1.7163, Val Acc: 62.88%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [100/100], Train Loss: 0.0181, Train Acc: 99.83%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [100/100], Val Loss: 1.7643, Val Acc: 62.46%
Model saved at epoch -> 100

```



	epoch	Train Acc	Train Loss	Val Acc	Val Loss
0	1	8.083538	3.291836	8.622503	3.245087
1	2	9.336609	3.185947	8.622503	3.198747
2	3	9.312039	3.153689	7.570978	3.225753
3	4	9.066339	3.127509	8.937960	3.111717
4	5	10.442260	3.109712	10.304942	3.107517
..
95	96	99.926290	0.005755	63.196635	1.715781
96	97	99.877150	0.009593	63.091483	1.830147
97	98	99.901720	0.011737	62.986330	1.718449
98	99	99.901720	0.007828	62.881178	1.716350
99	100	99.828010	0.018127	62.460568	1.764270

[100 rows x 5 columns]

Model saved at epoch 100

2. Best Loss (Mission3-2) Mission3-2 .

```
[ ]: #
LR = 0.001
BATCH_SIZE = 64
NUM_EPOCHS = 50
WEIGHT_DECAY = 5e-4

# ,
EXP_NAME = 'best_loss_exp'
SAVE_DIR = "Mission1"

#
train_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
val_transform = transforms.Compose([
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
data_loader = data_loading(
    tensor_dir='dataset/tensor/resize_tensor',
    batch_size=BATCH_SIZE,
    train_transform=train_transform,
    val_transform=val_transform
)

# ResNet18 ( X)
```

```

model = models.resnet18(weights=None, num_classes=31)
model = model.to(device)

#
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=LR, weight_decay=WEIGHT_DECAY)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=20, gamma=0.1)

#
training(
    experiment_name=EXP_NAME,
    save_dir=SAVE_DIR,
    data_loader=data_loader,
    optimizer=optimizer,
    criterion=criterion,
    scheduler=scheduler
)

```

Train Dataset: 4070, Val Dataset: 951

0%| | 0/50 [00:00<?, ?it/s]

0%| | 0/64 [00:00<?, ?it/s]

Epoch [1/50], Train Loss: 3.2918, Train Acc: 8.08%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [1/50], Val Loss: 3.2451, Val Acc: 8.62%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [2/50], Train Loss: 3.1859, Train Acc: 9.34%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [2/50], Val Loss: 3.1987, Val Acc: 8.62%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [3/50], Train Loss: 3.1537, Train Acc: 9.31%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [3/50], Val Loss: 3.2258, Val Acc: 7.57%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [4/50], Train Loss: 3.1275, Train Acc: 9.07%

0%| | 0/15 [00:00<?, ?it/s]

Epoch [4/50], Val Loss: 3.1117, Val Acc: 8.94%

0%| | 0/64 [00:00<?, ?it/s]

Epoch [5/50], Train Loss: 3.1097, Train Acc: 10.44%

```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [5/50], Val Loss: 3.1075, Val Acc: 10.30%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [6/50], Train Loss: 3.0818, Train Acc: 10.91%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [6/50], Val Loss: 3.1590, Val Acc: 8.62%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [7/50], Train Loss: 3.0633, Train Acc: 11.13%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [7/50], Val Loss: 3.2435, Val Acc: 8.73%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [8/50], Train Loss: 3.0566, Train Acc: 11.15%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [8/50], Val Loss: 3.0601, Val Acc: 10.83%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [9/50], Train Loss: 3.0382, Train Acc: 11.87%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [9/50], Val Loss: 3.0709, Val Acc: 9.04%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [10/50], Train Loss: 3.0268, Train Acc: 11.47%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [10/50], Val Loss: 3.1901, Val Acc: 7.57%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [11/50], Train Loss: 3.0122, Train Acc: 12.80%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [11/50], Val Loss: 3.0348, Val Acc: 11.25%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [12/50], Train Loss: 2.9758, Train Acc: 12.85%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [12/50], Val Loss: 3.1681, Val Acc: 11.78%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [13/50], Train Loss: 2.9698, Train Acc: 13.02%

```



```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [13/50], Val Loss: 3.1099, Val Acc: 9.57%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [14/50], Train Loss: 2.9270, Train Acc: 14.50%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [14/50], Val Loss: 3.2720, Val Acc: 8.62%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [15/50], Train Loss: 2.9132, Train Acc: 14.47%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [15/50], Val Loss: 2.9930, Val Acc: 12.41%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [16/50], Train Loss: 2.8835, Train Acc: 14.72%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [16/50], Val Loss: 3.4139, Val Acc: 11.46%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [17/50], Train Loss: 2.8584, Train Acc: 15.70%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [17/50], Val Loss: 3.1209, Val Acc: 11.78%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [18/50], Train Loss: 2.8288, Train Acc: 15.77%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [18/50], Val Loss: 2.9293, Val Acc: 14.09%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [19/50], Train Loss: 2.7936, Train Acc: 16.12%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [19/50], Val Loss: 2.8322, Val Acc: 16.82%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [20/50], Train Loss: 2.7467, Train Acc: 17.32%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [20/50], Val Loss: 3.2619, Val Acc: 12.72%
Model saved at epoch -> 20
0%|          | 0/64 [00:00<?, ?it/s]

```

Epoch [21/50], Train Loss: 2.5913, Train Acc: 21.65%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [21/50], Val Loss: 2.6374, Val Acc: 21.98%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [22/50], Train Loss: 2.4718, Train Acc: 24.91%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [22/50], Val Loss: 2.5904, Val Acc: 23.66%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [23/50], Train Loss: 2.3967, Train Acc: 26.81%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [23/50], Val Loss: 2.5349, Val Acc: 24.71%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [24/50], Train Loss: 2.3303, Train Acc: 28.60%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [24/50], Val Loss: 2.5111, Val Acc: 24.50%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [25/50], Train Loss: 2.2680, Train Acc: 30.96%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [25/50], Val Loss: 2.4444, Val Acc: 28.50%
Model saved at epoch -> 25

0%| | 0/64 [00:00<?, ?it/s]

Epoch [26/50], Train Loss: 2.1860, Train Acc: 33.19%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [26/50], Val Loss: 2.4172, Val Acc: 27.87%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [27/50], Train Loss: 2.0888, Train Acc: 36.88%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [27/50], Val Loss: 2.3565, Val Acc: 30.39%
0%| | 0/64 [00:00<?, ?it/s]

Epoch [28/50], Train Loss: 1.9911, Train Acc: 41.35%
0%| | 0/15 [00:00<?, ?it/s]

Epoch [28/50], Val Loss: 2.2817, Val Acc: 35.23%

```

0%|          | 0/64 [00:00<?, ?it/s]
Epoch [29/50], Train Loss: 1.8681, Train Acc: 43.76%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [29/50], Val Loss: 2.3246, Val Acc: 33.23%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [30/50], Train Loss: 1.7188, Train Acc: 50.07%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [30/50], Val Loss: 2.2529, Val Acc: 37.43%
Model saved at epoch -> 30
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [31/50], Train Loss: 1.5857, Train Acc: 55.41%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [31/50], Val Loss: 2.1496, Val Acc: 41.54%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [32/50], Train Loss: 1.4051, Train Acc: 62.38%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [32/50], Val Loss: 2.1051, Val Acc: 45.43%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [33/50], Train Loss: 1.2533, Train Acc: 68.01%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [33/50], Val Loss: 1.9971, Val Acc: 45.74%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [34/50], Train Loss: 1.0661, Train Acc: 75.01%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [34/50], Val Loss: 1.9793, Val Acc: 50.89%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [35/50], Train Loss: 0.8798, Train Acc: 81.84%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [35/50], Val Loss: 1.9858, Val Acc: 51.52%
Model saved at epoch -> 35
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [36/50], Train Loss: 0.7236, Train Acc: 86.54%
0%|          | 0/15 [00:00<?, ?it/s]

```

Epoch [36/50], Val Loss: 1.8308, Val Acc: 55.42%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [37/50], Train Loss: 0.5654, Train Acc: 91.28%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [37/50], Val Loss: 1.7918, Val Acc: 59.41%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [38/50], Train Loss: 0.4457, Train Acc: 94.08%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [38/50], Val Loss: 1.7235, Val Acc: 60.67%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [39/50], Train Loss: 0.3453, Train Acc: 96.68%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [39/50], Val Loss: 1.7110, Val Acc: 60.78%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [40/50], Train Loss: 0.2569, Train Acc: 97.81%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [40/50], Val Loss: 1.7076, Val Acc: 61.93%
 Model saved at epoch -> 40

0%| | 0/64 [00:00<?, ?it/s]

Epoch [41/50], Train Loss: 0.1794, Train Acc: 98.89%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [41/50], Val Loss: 1.5368, Val Acc: 62.15%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [42/50], Train Loss: 0.1389, Train Acc: 99.29%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [42/50], Val Loss: 1.5244, Val Acc: 62.67%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [43/50], Train Loss: 0.1298, Train Acc: 99.43%
 0%| | 0/15 [00:00<?, ?it/s]

Epoch [43/50], Val Loss: 1.5173, Val Acc: 62.67%
 0%| | 0/64 [00:00<?, ?it/s]

Epoch [44/50], Train Loss: 0.1197, Train Acc: 99.53%

```

0%|          | 0/15 [00:00<?, ?it/s]
Epoch [44/50], Val Loss: 1.5095, Val Acc: 63.20%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [45/50], Train Loss: 0.1066, Train Acc: 99.63%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [45/50], Val Loss: 1.5105, Val Acc: 63.20%
Model saved at epoch -> 45
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [46/50], Train Loss: 0.1056, Train Acc: 99.71%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [46/50], Val Loss: 1.5112, Val Acc: 62.67%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [47/50], Train Loss: 0.1021, Train Acc: 99.61%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [47/50], Val Loss: 1.5080, Val Acc: 62.88%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [48/50], Train Loss: 0.0953, Train Acc: 99.68%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [48/50], Val Loss: 1.5056, Val Acc: 62.99%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [49/50], Train Loss: 0.0909, Train Acc: 99.75%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [49/50], Val Loss: 1.5046, Val Acc: 63.09%
0%|          | 0/64 [00:00<?, ?it/s]
Epoch [50/50], Train Loss: 0.0878, Train Acc: 99.75%
0%|          | 0/15 [00:00<?, ?it/s]
Epoch [50/50], Val Loss: 1.5017, Val Acc: 63.20%
Model saved at epoch -> 50

```



	epoch	Train Acc	Train Loss	Val Acc	Val Loss
0	1	8.083538	3.291836	8.622503	3.245087
1	2	9.336609	3.185947	8.622503	3.198747
2	3	9.312039	3.153689	7.570978	3.225753
3	4	9.066339	3.127509	8.937960	3.111717
4	5	10.442260	3.109712	10.304942	3.107517
5	6	10.909091	3.081762	8.622503	3.158956
6	7	11.130221	3.063327	8.727655	3.243517
7	8	11.154791	3.056645	10.830705	3.060093
8	9	11.867322	3.038169	9.043113	3.070874
9	10	11.474201	3.026791	7.570978	3.190090
10	11	12.800983	3.012236	11.251314	3.034799
11	12	12.850123	2.975829	11.777077	3.168130
12	13	13.022113	2.969792	9.568875	3.109891
13	14	14.496314	2.927012	8.622503	3.271963
14	15	14.471744	2.913243	12.407992	2.992993
15	16	14.717445	2.883492	11.461619	3.413853
16	17	15.700246	2.858416	11.777077	3.120889
17	18	15.773956	2.828796	14.090431	2.929308
18	19	16.117936	2.793589	16.824395	2.832158
19	20	17.321867	2.746675	12.723449	3.261891
20	21	21.646192	2.591321	21.976866	2.637435
21	22	24.914005	2.471822	23.659306	2.590398
22	23	26.805897	2.396749	24.710831	2.534911
23	24	28.599509	2.330283	24.500526	2.511128
24	25	30.958231	2.267968	28.496320	2.444427
25	26	33.194103	2.186004	27.865405	2.417221
26	27	36.879607	2.088790	30.389064	2.356525
27	28	41.351351	1.991087	35.226078	2.281652

28	29	43.759214	1.868078	33.228181	2.324558
29	30	50.073710	1.718838	37.434280	2.252931
30	31	55.405405	1.585710	41.535226	2.149633
31	32	62.383292	1.405115	45.425868	2.105072
32	33	68.009828	1.253278	45.741325	1.997073
33	34	75.012285	1.066142	50.893796	1.979315
34	35	81.842752	0.879833	51.524711	1.985802
35	36	86.535627	0.723647	55.415352	1.830799
36	37	91.277641	0.565436	59.411146	1.791835
37	38	94.078624	0.445692	60.672976	1.723513
38	39	96.683047	0.345340	60.778128	1.711010
39	40	97.813268	0.256905	61.934805	1.707625
40	41	98.894349	0.179391	62.145110	1.536823
41	42	99.287469	0.138889	62.670873	1.524394
42	43	99.434889	0.129757	62.670873	1.517275
43	44	99.533170	0.119672	63.196635	1.509525
44	45	99.631450	0.106621	63.196635	1.510459
45	46	99.705160	0.105647	62.670873	1.511184
46	47	99.606880	0.102098	62.881178	1.507973
47	48	99.680590	0.095299	62.986330	1.505559
48	49	99.754300	0.090932	63.091483	1.504602
49	50	99.754300	0.087774	63.196635	1.501733

Model saved at epoch 50

1.2.6 5.

```
[ ]: # best acc model

SAVE_DIR = "Mission1"
EXP_NAME = "best_acc_exp"

best_path = f"./{SAVE_DIR}/{EXP_NAME}/best_model.pth"
model.load_state_dict(torch.load(best_path, weights_only=True))

def evaluate_model(model, data_loader):
    model.eval()
    val_loss = 0
    val_correct = 0
    total = 0
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for images, labels in tqdm(data_loader, leave=False):
            images, labels = images.to(device), labels.to(device)

            outputs = model(images)
```

```

        loss = criterion(outputs, labels)
        val_loss += loss.item()

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        val_correct += (predicted == labels).sum().item()

        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

    val_loss /= len(data_loader)
    val_acc = 100 * val_correct / total

    print(f"Best Model Performance | Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")

    return np.array(all_preds), np.array(all_labels), val_loss, val_acc

y_pred, y_true, val_loss, val_acc = evaluate_model(model, data_loader[1])

```

0%| | 0/15 [00:00<?, ?it/s]

Best Model Performance | Val Loss: 1.8487, Val Acc: 63.83%

1.2.7 6.

1. .

```

[ ]: # class
import json
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report

#
def plot_confusion_matrix(y_true, y_pred, classes):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(20, 16))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=classes,
        yticklabels=classes)
    plt.title('Confusion Matrix')
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.savefig(f'{SAVE_DIR}/{EXP_NAME}/pred_heatmap.png', bbox_inches='tight')
    plt.show()

#
def print_classification_report(y_true, y_pred, classes):

```




```

report = classification_report(y_true, y_pred, target_names=classes,
↪digits=3)
print(report)

#
label_map_ev = json.load(open('./dataset/label_map.json'))

#
class_names = [label_map_ev[str(i)] for i in range(31)]
plot_confusion_matrix(y_true, y_pred, class_names)
print_classification_report(y_true, y_pred, class_names)

```



	precision	recall	f1-score	support
W_minimal	0.700	0.600	0.646	35
M_sportivecasual	0.538	0.538	0.538	52

W_normcore	0.179	0.250	0.208	20
M_bold	0.651	0.491	0.560	57
M_normcore	0.288	0.588	0.387	51
W_classic	0.750	0.682	0.714	22
W_feminine	0.623	0.750	0.680	44
W_sportivecasual	0.800	0.750	0.774	48
M_hippie	0.773	0.707	0.739	82
M_metrosexual	0.580	0.690	0.630	58
W_ecology	0.588	0.588	0.588	17
M_hiphop	0.800	0.667	0.727	66
W_hippie	0.750	0.643	0.692	14
M_ivy	0.740	0.722	0.731	79
W_bodyconscious	0.750	0.652	0.698	23
W_powersuit	0.885	0.676	0.767	34
W_kitsch	0.778	0.636	0.700	22
W_disco	0.571	0.400	0.471	10
M_mods	0.770	0.713	0.740	80
W_cityglam	0.476	0.556	0.513	18
W_hiphop	0.400	0.500	0.444	8
W_punk	0.556	0.417	0.476	12
W_military	0.714	0.556	0.625	9
W_lounge	0.083	0.125	0.100	8
W_oriental	0.923	0.667	0.774	18
W_athleisure	0.643	0.643	0.643	14
W_genderless	0.900	0.750	0.818	12
W_popart	0.857	0.750	0.800	8
W_lingerie	0.800	0.800	0.800	5
W_space	1.000	0.600	0.750	15
W_grunge	0.667	0.600	0.632	10
accuracy			0.638	951
macro avg	0.662	0.603	0.625	951
weighted avg	0.679	0.638	0.651	951

```
[ ]: # 10
class_accuracy = classification_report(y_true, y_pred,
    ↪target_names=class_names, output_dict=True)
class_accuracy = {k: v['f1-score'] for k, v in class_accuracy.items() if k in
    ↪class_names}
worst_classes = sorted(class_accuracy.items(), key=lambda x: x[1])

print("Worst performing classes:")
for i, (cls, score) in enumerate(worst_classes):
    print(f"{cls}: {score:.3f}")
    if i == 9:
        print("=====")
```

Worst performing classes:

W_lounge: 0.100

W_normcore: 0.208

M_normcore: 0.387

W_hiphop: 0.444

W_disco: 0.471

W_punk: 0.476

W_cityglam: 0.513

M_sportivecasual: 0.538

M_bold: 0.560

W_ecology: 0.588

=====

W_military: 0.625

M_metrosexual: 0.630

W_grunge: 0.632

W_athleisure: 0.643

W_minimal: 0.646

W_feminine: 0.680

W_hippie: 0.692

W_bodyconscious: 0.698

W_kitsch: 0.700

W_classic: 0.714

M_hiphop: 0.727

M_ivy: 0.731

M_hippie: 0.739

M_mods: 0.740

W_space: 0.750

W_powersuit: 0.767

W_sportivecasual: 0.774

W_oriental: 0.774

W_popart: 0.800

W_lingerie: 0.800

W_genderless: 0.818